

Conceptualized Representation Learning for Chinese Biomedical Text Mining

Ningyu Zhang
ningyu.zny@alibaba-inc.com
Alibaba Group
China

Liang Dong
liang.dongl1@alibaba-inc.com
Alibaba Group
China

Qianghuai Jia
qianghuai.jqh@alibaba-inc.com
Alibaba Group
China

Feng Gao
gf142364@alibaba-inc.com
Alibaba Group
China

Kangping Yin
kangping.yinkp@alibaba-inc.com
Alibaba Group
China

Nengwei Hua
nengwei.huanw@alibaba-inc.com
Alibaba Group
China

ABSTRACT

Biomedical text mining is becoming increasingly important as the number of biomedical documents and web data rapidly grows. Recently, word representation models such as BERT has gained popularity among researchers. However, it is difficult to estimate their performance on datasets containing biomedical texts as the word distributions of general and biomedical corpora are quite different. Moreover, the medical domain has long-tail concepts and terminologies that are difficult to be learned via language models. For the Chinese biomedical text, it is more difficult due to its complex structure and the variety of phrase combinations. In this paper, we investigate how the recently introduced pre-trained language model BERT can be adapted for Chinese biomedical corpora and propose a novel conceptualized representation learning approach. We also release a new Chinese Biomedical Language Understanding Evaluation benchmark (**ChineseBLUE**). We examine the effectiveness of Chinese pre-trained models: BERT, BERT-wm, RoBERTa, and our approach. Experimental results on the benchmark show that our approach could bring significant gain. We release the pre-trained model on GitHub: <https://github.com/alibaba-research/ChineseBLUE>.

CCS CONCEPTS

• **Information systems** → **Information extraction**.

KEYWORDS

Chinese Biomedical Natural Language Processing; Conceptualized Representation Learning; BERT

ACM Reference Format:

Ningyu Zhang, Qianghuai Jia, Kangping Yin, Liang Dong, Feng Gao, and Nengwei Hua. 2020. Conceptualized Representation Learning for Chinese Biomedical Text Mining. In *WSDM '20*, February 3–7, 2020, Houston, ACM, New York, NY, USA, 4 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

WSDM '20, February 3–7, 2020, Houston

© 2020 Association for Computing Machinery.
ACM ISBN 978-x-xxxx-xxxx-x/YY/MM...\$15.00
<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 INTRODUCTION

Nowadays, the volume of biomedical literature and biomedical web pages continues to increase rapidly. Lots of new articles and web pages containing biomedical discoveries and new insights are continuously published. Indeed, there is an increasingly high demand for biomedical text mining. Recent progress in the biomedical text mining approach is made possible by the development of deep learning techniques used in natural language processing (NLP). For example, pre-trained language models such as BERT [2], ERNIE [6], XLNet [7] and RoBERTa [5] have demonstrated remarkable successes in modeling contextualized word representations by utilizing the massive amount of training text. As a fundamental technique in natural language processing (NLP), the language models pre-trained on text could be easily transferred to learn downstream NLP tasks with finetuning, which achieve the state-of-the-art performances on many tasks including named entity recognition, paraphrase identification, question answering and information retrieval.

However, it has limitations to apply state-of-the-art NLP methodologies to biomedical text mining directly. Firstly, since recent representation models such as BERT are trained and tested mainly on general domain datasets such as Wikipedia, it is difficult to adapt to biomedical datasets without losing the performance. Moreover, the word distributions of general and biomedical text are quite different, which can be a problem for biomedical text mining. In addition, there exist long-tail concepts and terminologies in biomedical texts which are difficult to be learned via language models. For the Chinese biomedical text, it is somewhat more difficult due to its complex structure and the variety of phrase combinations. To this end, recent biomedical text mining models rely mostly on adapted versions of word representations [3]. Considering whether it is possible to automatically inject biomedical knowledge to the language representation learning for Chinese medical corpus, we hypothesize that current state-of-the-art word representation models such as BERT should be trained on biomedical corpora with prior biomedical knowledge to be effective in biomedical text mining tasks. However, there exist two problems: (1) how to retrieve the biomedical domain knowledge; (2) how to leverage such knowledge to the representation learning.

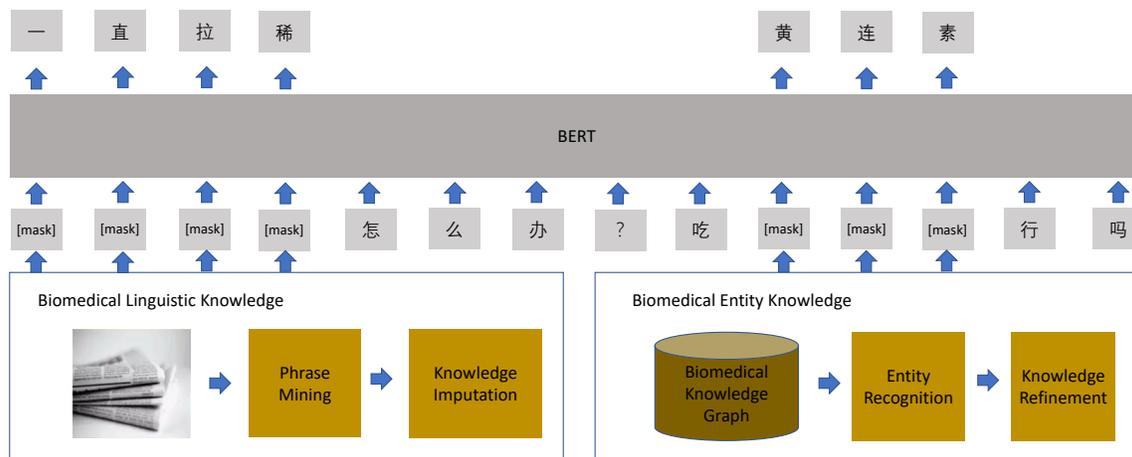


Figure 1: The structure of the conceptualized representation learning (C-BERT) model. We utilize whole entity masking which masks the entity “黄连素” (“berberine”) and whole span masking which masks the phrase “一直拉稀” (“often diarrhea”) to explicitly inject biomedical knowledge.

In this paper, we propose a conceptualize representation learning approach (MC-BERT) for Chinese biomedical language understanding. Specifically, we propose coarse-to-fine masking strategies to inject entity and linguistic domain knowledge into representation learning. As there are no benchmarks for the Chinese Biomedical Language Understanding Evaluation, we release the first large scale benchmark¹ including name entity recognition, paraphrase identification, question answering, information retrieval, intent detection, and text classification. Experiments show that our approach achieves state-of-art results.

2 RELATED WORK

It is effective to learn general language representation by pre-training a language model with a large amount of unannotated data. Recently, several studies BERT [1], BERT-wwm [1], RoBERTa [5], XLNet [7] which centered on contextualized language representations have been proposed and context-dependent language representations have shown state-of-the-art results in various natural language processing tasks. Moreover, recent studies have shown that injecting extra knowledge information can significantly enhance original models, such as knowledge acquisition [8]. Some work has attempted to joint representation learning of words and entities for effectively leveraging external Knowledge Graphs and achieved promising results [9]. [6] propose the knowledge masking strategy for masked language models to enhance language representation by knowledge. However, there is only a little research on pretraining in the biomedical domain. [4] firstly proposed BioBERT, which is a domain-specific language representation model pre-trained on large-scale biomedical corpora. In this paper, we further leverage both corpora and knowledge graphs to train an enhanced language representation model based on BERT. To the best of our knowledge, we are the first approach to inject biomedical knowledge into Chinese biomedical representation leaning.

¹<https://github.com/alibaba-research/ChineseBLUE>

3 APPROACH

We present MC-BERT, a self-supervised pretraining method designed to represent the text better. Our approach is inspired by BERT but deviates from its bi-text classification framework in three ways. First, we use a different mask generation procedure to mask spans of tokens, rather than only random ones. We also introduce two kinds of masking strategies, namely whole entity masking and whole span masking. Finally, MC-BERT split the input document into segments based on the actual “sentences” provided by the user as positive samples and sample random sentences from other documents as negative samples for the next sentence prediction.

3.1 Whole Entity Masking

Different from whole word masking, we do not mask random words but medical entities such as “腹痛” (“stomach ache”), which can explicitly inject medical domain knowledge. We utilize the Chinese biomedical knowledge graph and biomedical named entity recognition to extract and refine entities in the medical domains, including syndrome, decease, examination, treatment, drug, and so on.

3.2 Whole Span Masking

Due to the complex structure and the variety of phrase combinations in Chinese, we also inject fine-grained biomedical knowledge with whole span masking. For example, the phrases “肚子有一点疼” (“a little pain in the stomach”), “腹部一阵一阵痛” (“a pain in the abdomen”), “腹痛” (“stomach ache”) have the same meaning with concepts “腹痛” (“stomach ache”), the whole entity masking cannot explicitly inject such knowledge. We argue that the prior linguistic knowledge of phrases is necessary for Chinese biomedical language understanding. Specifically, we firstly utilize Autphrase² to extract phrases. We also retrieve common biomedical phrases from Alibaba Cognitive Concept Graph³. Then

²<https://github.com/shangjingbo1226/AutoPhrase>

³<https://github.com/alibaba-research/CognitiveConceptGraph>

Table 1: Statistics of pretraining corpus.

Corpus Type	#Sentences
Chinese Biomedical Community QA	20M
Chinese Medical Encyclopedia	100K
Chinese Electric Medical Record	10K

we leverage domain rules to augment data and train a binary classifier to filter those none-biomedical phrases. We collect the n-gram of entities and attribute words in the medical encyclopedia as positive samples and generate random phrases as negative samples for the classifier via fastText⁴.

Note that all of the words in the same unit are masked during word representation training instead of only one word or character being masked. In this way, the prior knowledge of phrases and entities are explicitly learned during the training procedure. Instead of adding the knowledge embedding directly, MC-BERT explicitly learned the information about knowledge to guide word embedding learning. This can make the model have better generalization and adaptability in the biomedical domain.

3.3 Further Pretraining in Biomedical Domain

We did not train our model from scratch but from the BERT-base. We trained 100K steps on the samples with a maximum length of 512, with an initial learning rate of $1e-5^5$. Note that the learning rate is critical, and we **do not** use the learning rate warmup as we empirically find it will lead to severe catastrophe forgetting in the biomedical domain. The overall training procedure is shown below.

Algorithm 1 Overall Training Procedure of MC-BERT

- 1: Generate candidate biomedical entities and refine them from the biomedical knowledge graph.
 - 2: Generate candidate phrases via Autophrase from the raw corpus.
 - 3: Augment and filter those phrases via rules and fastText.
 - 4: Duplicate and shuffle corpus ten times and generate whole entity/span masking training samples with rate 15%.
 - 5: Initialize all parameters with BERT-base and make further pre-training in the biomedical corpus.
-

4 EXPERIMENTS

4.1 Pretraining Data and Settings

For the Chinese corpus, we collect a variety of data, such as Chinese community biomedical question answering, Chinese medical encyclopedia, Chinese electronic health records (EHR), and so on from Alibaba Shenma Search Engine⁶. The details of the pretraining corpus are shown in Table 1.

To compare with BERT, we leverage the same model settings of the transformer as BERT. The base model contains 12 layers, 12

Table 2: Statistics of ChineseGLUE.

Dataset	Train	Dev	Test	Task	Metric
cEHRNER	914	44	41	NER	F1
cMedQANER	1,673	175	215	NER	F1
cMedQQ	16,071	1,793	1,935	PI	F1
cMedQNLI	80,950	9,065	9,969	QA	F1
cMedQA	49,719	5,475	6,149	QA	F1
cMedIR	80,000	10,000	10,000	IR	PAIR
cMedIC	1,683	123	84	IC	F1
cMedTC	14,610	1,550	1,800	TC	F1

self-attention heads, and 768-dimensional of hidden size, while the large model contains 24 layers, 16 self-attention heads, and 1024-dimensional of hidden size.

4.2 Finetuning Tasks

We executed extensive experiments on Chinese NLP tasks and release a Chinese Biomedical Language Understanding Evaluation benchmark (ChineseBLUE), as shown in table 2. The following Chinese datasets in the ChineseBLUE are chosen to evaluate the performance of MC-BERT on Chinese tasks:

- **Named Entity Recognition (NER)** aims to recognize various entities, including diseases, drugs, syndromes, etc. The cEHRNER dataset labeled from the Chinese electronic health records and the cMedQANER dataset labeled from Chinese community question answering is chosen.
- **Paraphrase Identification (PI)** aims to identify whether two sentences express the same meaning. We use cMedQQ, which consists of search query pairs.
- **Question Answering (QA)**, which can be approximated as ranking candidate answer sentences based on their similarity. We assign 0,1 labels to the QA pairs, which convert to the binary classification problem. We use cMedQNLI, which consists of long answers and cMeQA, which consists of short answers.
- **Information Retrieval (IR)**, which aims to retrieve most related documents given search queries. IR can be regarded as a ranking task. We adopt the **PAIR**⁷ score to evaluate the model. We use the cMedIR dataset, which consists of queries with multiple documents and their relative scores.
- **Intent Classification (IC)** aims to assign intent labels to the queries, which can be regarded as multiple label classification tasks. We use the cMedIC dataset, which consists of queries with three intent labels (e.g., no intention, weak intention, and firm intention).
- **Text Classification (TC)** aims to assign multiple labels to the sentence. We use the cMedTC dataset, which consists of biomedical texts with multiple labels.

⁴<https://github.com/facebookresearch/fastText>

⁵The initial learning rate should be adapted for different corpus

⁶<http://m.sm.cn>

⁷A popular NDCG-like ranking metric in the search engine, which refers to the number of positive ranked documents divide the number of negative ranked documents.

Table 3: BERT and MC-BERT results on different tasks of ChineseBLUE.

Model	cEHRNER	cMedQANER	cMedQQ	cMedQA
BERT-base	88.2	86.3	86.5	81.0
MC-BERT	90.0	88.1	87.5	82.3
Model	cMedQNLI	cMedIR	cMedIC	cMedTC
BERT-base	93.3	1.77	86.0	79.0
MC-BERT	95.5	2.04	87.5	82.1

Table 4: NER Results of different Chinese pretraining models and ablation study.

Model	Precision	Rrecall	F1
C-BERT	89.8	90.4	90.0
w/o entity	89.5	89.6	89.6
w/o span	88.1	88.3	88.2
BERT-base	88.0	88.4	88.2
BERT-wwm	89.1	89.2	89.2
RoBERTa	89.1	89.5	89.3

4.3 Results

Table 3 shows the performances of classical Chinese NLP tasks. It can be seen that MC-BERT outperforms BERT-base on all tasks, including cEHRNER, cMedQANER, cMedQQ, cMedQA, yet the performance achieve only a little progress on the rest, which is caused by the difference in pretraining between the two methods. Specifically, the pretraining data of MC-BERT does not contain instances whose length less than 128, but the length of finetuning instances in datasets such as cMedTC is less than 128. Specifically, MC-BERT yields improvements of more than 2 points over BERT-base on the name entity recognition, and yields improvements of more than 1 point over BERT-base on the paraphrase identification task.

4.4 Analysis of Chinese Pretrained Models

To better demonstrate the performance of different strategies in our model and different pretraining approaches in Chinese, we separately remove the whole entity and span masking and also compare our model with other Chinese pretraining models. The experimental results of named entity recognition on the cEHRNER dataset are summarized in Table 4. **C-BERT** is our method; **w/o entity** is the method without whole entity masking; **w/o span** is the method without whole span masking; **BERT-wwm** [1] is a whole word masking pretraining approach for the Chinese text; **RoBERTa**⁸ [5] is a robustly optimized BERT pretraining approach. We observe that: (1) Performance degrades when we remove "whole entity mask" and "whole span mask." This is reasonable because the entity and span explicitly inject domain knowledge into representation learning. (2) Our approach performs better than the recent state-of-art pretraining approach RoBERTa in Chinese, which indicates further pretraining in the biomedical domain is necessary.

⁸We adopt all the pretrained Chinese model from <https://github.com/ymcui/Chinese-BERT-wwm>

4.5 Discussion

As far as we concern, there are only a few Chinese open biomedical datasets that hinder the research of Chinese biomedical language understanding. We make the first step to build a Chinese Biomedical language understanding benchmark and investigate the pre-training strategies. In our experiments, we notice that there exist long-tail terminologies in Chinese biomedical corpus such as "肌酶" ("muscle enzyme"). The vanilla mask language model can not learn robust representations from only word-level masking. Therefore, we propose the whole entity/span masking to inject domain knowledge explicitly, which improves the performance. Note that our approach is language-agnostic, it is convenient to adapt our approach to other languages such as English and even low-resource languages. However, our work has some limitations. First, our work relies on phrase mining, which makes the whole training procedure not end-to-end and may result in error propagation. Otherwise, there is also some concept knowledge in the biomedical domain, such as various phrases refer to the same normalized concept. We believe the performance of our model can be further improved by injecting such knowledge, which will be part of our future work.

5 CONCLUSION AND FUTURE WORK

In this article, we introduce MC-BERT, which is a pre-trained language representation model with domain knowledge for biomedical text mining. We show that pretraining BERT on biomedical corpora is crucial in applying it to the biomedical domain. The ChineseBLUE and MC-BERT will soon be available to the BioNLP community. Our motivation is to develop a universal, GLUE-like, and open platform for the Chinese BioNLP community and a composable and generalized representation algorithm to inject domain knowledge. Our work is but a small step in this direction.

REFERENCES

- [1] Yiming Cui, Wanxiang Che, Ting Liu, Bing Qin, Ziqing Yang, Shijin Wang, and Guoping Hu. 2019. Pre-Training with Whole Word Masking for Chinese BERT. *arXiv preprint arXiv:1906.08101* (2019).
- [2] Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2018. Bert: Pre-training of deep bidirectional transformers for language understanding. *arXiv preprint arXiv:1810.04805* (2018).
- [3] Maryam Habibi, Leon Weber, Mariana Neves, David Luis Wiegandt, and Ulf Leser. 2017. Deep learning with word embeddings improves biomedical named entity recognition. *Bioinformatics* 33, 14 (2017), i37–i48.
- [4] Jinhyuk Lee, Wonjin Yoon, Sungdong Kim, Donghyeon Kim, Sunkyu Kim, Chan Ho So, and Jaewoo Kang. 2019. Biobert: pre-trained biomedical language representation model for biomedical text mining. *arXiv preprint arXiv:1901.08746* (2019).
- [5] Yinhan Liu, Myle Ott, Naman Goyal, Jingfei Du, Mandar Joshi, Danqi Chen, Omer Levy, Mike Lewis, Luke Zettlemoyer, and Veselin Stoyanov. 2019. Roberta: A robustly optimized bert pretraining approach. *arXiv preprint arXiv:1907.11692* (2019).
- [6] Yu Sun, Shuohuan Wang, Yukun Li, Shikun Feng, Xuyi Chen, Han Zhang, Xin Tian, Danxiang Zhu, Hao Tian, and Hua Wu. 2019. ERNIE: Enhanced Representation through Knowledge Integration. *arXiv preprint arXiv:1904.09223* (2019).
- [7] Zhilin Yang, Zihang Dai, Yiming Yang, Jaime Carbonell, Ruslan Salakhutdinov, and Quoc V Le. 2019. XLNet: Generalized Autoregressive Pretraining for Language Understanding. *arXiv preprint arXiv:1906.08237* (2019).
- [8] Ningyu Zhang, Shumin Deng, Zhanlin Sun, Guanying Wang, Xi Chen, et al. 2019. Long-tail Relation Extraction via Knowledge Graph Embeddings and Graph Convolution Networks. In *In NAACL*.
- [9] Ningyu Zhang, Zhanlin Sun, Shumin Deng, Jiaoyan Chen, and Huajun Chen. 2019. Improving Few-shot Text Classification via Pretrained Language Representations. *arXiv preprint arXiv:1908.08788* (2019).